

```
#first load all the important library is needed for prediction
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
import plotly.express as px
import plotly.io as pio
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
#collect the data from the file as csv data
df=pd.read_csv("/content/mobile_data (1).csv")
```

```
#load the mobile dataset by using read_csv function which load the csv file and it make the data tabular form
df
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	75
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	198
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	171
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	178
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	121
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	189
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	196
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	163
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	67
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	75

2000 rows × 21 columns



```
#perform EDA analysis
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   battery_power         2000 non-null  int64
1   blue                  2000 non-null  int64
2   clock_speed           2000 non-null  float64
3   dual_sim              2000 non-null  int64
4   fc                    2000 non-null  int64
5   four_g                2000 non-null  int64
6   int_memory            2000 non-null  int64
7   m_dep                 2000 non-null  float64
8   mobile_wt             2000 non-null  int64
9   n_cores               2000 non-null  int64
10  pc                    2000 non-null  int64
11  px_height              2000 non-null  int64
12  px_width               2000 non-null  int64
13  ram                   2000 non-null  int64
14  sc_h                  2000 non-null  int64
15  sc_w                  2000 non-null  int64
16  talk_time              2000 non-null  int64
```

```
17 three_g      2000 non-null   int64
18 touch_screen  2000 non-null   int64
19 wifi          2000 non-null   int64
20 price_range   2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

df.describe()

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.046500	0.501750	140.249000	...
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.145715	0.288416	35.399655	...
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000	0.100000	80.000000	...
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000	0.200000	109.000000	...
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000	0.500000	141.000000	...
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000	0.800000	170.000000	...
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000	1.000000	200.000000	...

8 rows × 21 columns



df.corr()

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_core
battery_power	1.000000	0.011252	0.011482	-0.041847	0.033334	0.015665	-0.004004	0.034085	0.001844	-0.029727
blue	0.011252	1.000000	0.021419	0.035198	0.003593	0.013443	0.041177	0.004049	-0.008605	0.036161
clock_speed	0.011482	0.021419	1.000000	-0.001315	-0.000434	-0.043073	0.006545	-0.014364	0.012350	-0.005724
dual_sim	-0.041847	0.035198	-0.001315	1.000000	-0.029123	0.003187	-0.015679	-0.022142	-0.008979	-0.024651
fc	0.033334	0.003593	-0.000434	-0.029123	1.000000	-0.016560	-0.029133	-0.001791	0.023618	-0.013351
four_g	0.015665	0.013443	-0.043073	0.003187	-0.016560	1.000000	0.008690	-0.001823	-0.016537	-0.029700
int_memory	-0.004004	0.041177	0.006545	-0.015679	-0.029133	0.008690	1.000000	0.006886	-0.034214	-0.028310
m_dep	0.034085	0.004049	-0.014364	-0.022142	-0.001791	-0.001823	0.006886	1.000000	0.021756	-0.003500
mobile_wt	0.001844	-0.008605	0.012350	-0.008979	0.023618	-0.016537	-0.034214	0.021756	1.000000	-0.018981
n_cores	-0.029727	0.036161	-0.005724	-0.024658	-0.013356	-0.029706	-0.028310	-0.003504	-0.018989	1.000000
pc	0.031441	-0.009952	-0.005245	-0.017143	0.644595	-0.005598	-0.033273	0.026282	0.018844	-0.001191
px_height	0.014901	-0.006872	-0.014523	-0.020875	-0.009990	-0.019236	0.010441	0.025263	0.000939	-0.006871
px_width	-0.008402	-0.041533	-0.009476	0.014291	-0.005176	0.007448	-0.008335	0.023566	0.000090	0.024480
ram	-0.000653	0.026351	0.003443	0.041072	0.015099	0.007313	0.032813	-0.009434	-0.002581	0.004861
sc_h	-0.029959	-0.002952	-0.029078	-0.011949	-0.011014	0.027166	0.037771	-0.025348	-0.033855	-0.000311
sc_w	-0.021421	0.000613	-0.007378	-0.016666	-0.012373	0.037005	0.011731	-0.018388	-0.020761	0.025820
talk_time	0.052510	0.013934	-0.011432	-0.039404	-0.006829	-0.046628	-0.002790	0.017003	0.006209	0.013141
three_g	0.011522	-0.030236	-0.046433	-0.014008	0.001793	0.584246	-0.009366	-0.012065	0.001551	-0.014731
touch_screen	-0.010516	0.010061	0.019756	-0.017117	-0.014828	0.016758	-0.026999	-0.002638	-0.014368	0.023771
wifi	-0.008343	-0.021863	-0.024471	0.022740	0.020085	-0.017620	0.006993	-0.028353	-0.000409	-0.009961
price_range	0.200723	0.020573	-0.006606	0.017444	0.021998	0.014772	0.044435	0.000853	-0.030302	0.004391

21 rows × 21 columns



df.isnull().sum()

```

battery_power    0
blue             0
clock_speed      0
dual_sim         0
fc               0
four_g           0
int_memory       0
m_dep            0
mobile_wt        0
n_cores          0
pc               0
px_height        0
px_width         0
ram              0
sc_h             0
sc_w             0
talk_time        0
three_g          0
touch_screen     0
wifi             0
price_range      0
dtype: int64

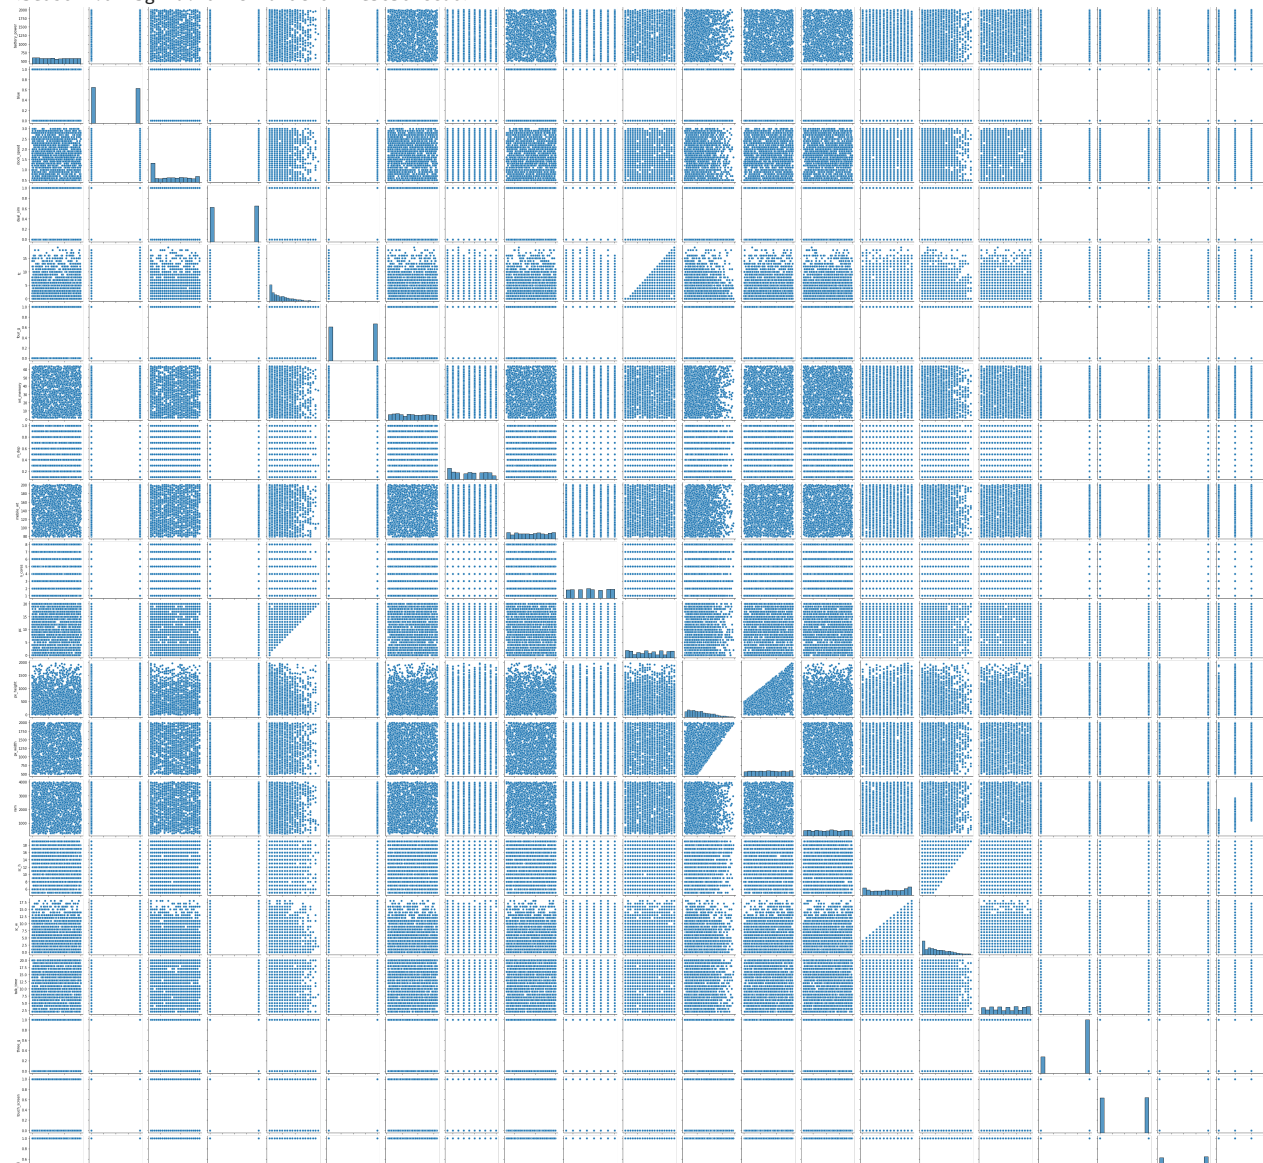
```

```
#visualization of different features by using different graphs representation
```

```
# we use the seaborn.pair plot for visualization weather there is linear or non linear model or the data having more corration
```

```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7fe3650fccd0>
```




```
# we will see the visualization of battery vs price_range
```

```
# Violin plot
```

```
sns.violinplot(x='price_range', y='px_height', data=df)
plt.xticks(rotation=90)
```


```
(array([0, 1, 2, 3]), <a list of 4 Text major ticklabel objects>)
```



```
#we use the plotly library for interactive visualization
```


```
ax= px.histogram(df,x= "battery_power", template= "plotly_dark",color= "px_height",title='battery power of mobile distribution')
ax.show()
```

battery power of mobile distribution




```
ax= px.histogram(df,x= "dual_sim", template= "plotly_dark",color= "sc_h",title='dual sim and screen height of mobile distribution')
ax.show()
```

dual sim and screen height of mobile distribution




distribution



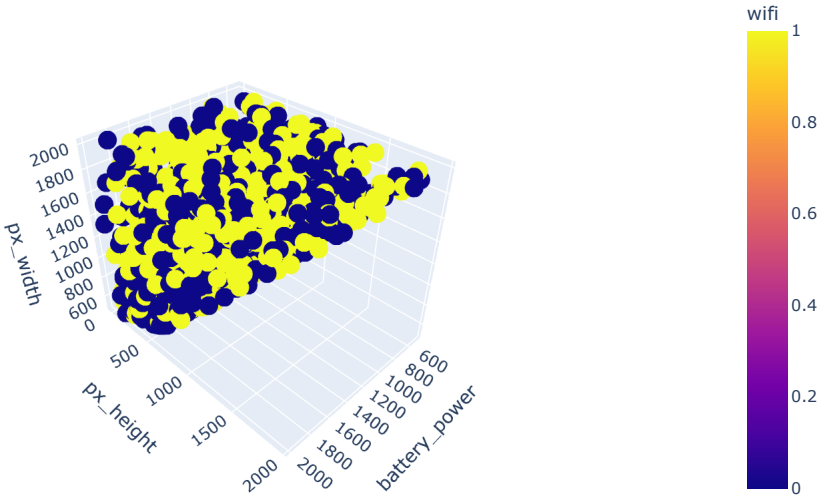
```
ax= px.pie(df, names= "sc_h",template= "plotly_dark",title= "specification of mobile phone",hole= 0.5)
ax.show()
```

specification of mobile phone



```
#scatter plot
```

```
fig = px.scatter_3d(df, x='battery_power', y='px_height', z='px_width',
                    color='wifi')
fig.show()
```



```
#we import plotly for better presentation
```

```
#sepaerating x and y for taining and testing
```

```
x = df.iloc[:, :-1]
y = df["price_range"]
```

x

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width
0	842	0	2.2	0	1	0	7	0.6	188	2	2	20	756
1	1021	1	0.5	1	0	1	53	0.7	136	3	6	905	1988
2	563	1	0.5	1	2	1	41	0.9	145	5	6	1263	1716
3	615	1	2.5	0	0	0	10	0.8	131	6	9	1216	1786
4	1821	1	1.2	0	13	1	44	0.6	141	2	14	1208	1212
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	14	1222	1890
1996	1965	1	2.6	1	0	0	39	0.2	187	4	3	915	1965
1997	1911	0	0.9	1	1	1	36	0.7	108	8	3	868	1632
1998	1512	0	0.9	0	4	1	46	0.1	145	5	5	336	670
1999	510	1	2.0	1	5	1	45	0.9	168	6	16	483	754

2000 rows × 20 columns



y

0	1
1	2

```

2      2
3      2
4      1
..
1995   0
1996   2
1997   3
1998   0
1999   3
Name: price_range, Length: 2000, dtype: int64

#split the dats x and y for training and testing

xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.3, random_state=1)

#import library of linear regression

# Fitting Multiple Linear Regression to the Training set
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(xtrain, ytrain)

LinearRegression()

# Predicting the Test set results
ypred = regressor.predict(xtest)

#evaluate the model
from sklearn.metrics import r2_score

r2_score(ytest,ypred)

0.909367881466235

#loss function
from sklearn.metrics import mean_absolute_error

mean_absolute_error(ytest,ypred)    #MAE

0.2751015081723019

from sklearn.metrics import mean_squared_error    #MSE

mean_squared_error(ytest,ypred)

0.10794260141782923

np.sqrt(mean_squared_error(ytest,ypred))

0.32854619373511124

#building different models by using classification ml algorithm
logreg = LogisticRegression()
knn = KNeighborsClassifier()
svm = SVC()
dt=DecisionTreeClassifier()

#using define function to create the fifferent model

def mymodel(model):
    model.fit(xtrain, ytrain)
    ypred = model.predict(xtest)
    print(classification_report(ytest, ypred))
    return model

mymodel(logreg)

```

	precision	recall	f1-score	support
0	0.78	0.80	0.79	135
1	0.53	0.58	0.55	149
2	0.53	0.40	0.45	168
3	0.67	0.78	0.72	148
accuracy			0.63	600
macro avg	0.63	0.64	0.63	600
weighted avg	0.62	0.63	0.62	600

```
LogisticRegression()
```

```
mymodel(knn)
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	135
1	0.91	0.91	0.91	149
2	0.88	0.89	0.88	168
3	0.94	0.91	0.93	148
accuracy			0.92	600
macro avg	0.92	0.92	0.92	600
weighted avg	0.92	0.92	0.92	600

```
KNeighborsClassifier()
```

```
mymodel(svm)
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	135
1	0.93	0.93	0.93	149
2	0.94	0.90	0.92	168
3	0.95	0.97	0.96	148
accuracy			0.94	600
macro avg	0.94	0.95	0.94	600
weighted avg	0.94	0.94	0.94	600

```
SVC()
```

```
mymodel(dt)
```

	precision	recall	f1-score	support
0	0.91	0.86	0.89	135
1	0.77	0.85	0.81	149
2	0.84	0.80	0.82	168
3	0.91	0.91	0.91	148
accuracy			0.85	600
macro avg	0.86	0.86	0.86	600
weighted avg	0.86	0.85	0.85	600

```
DecisionTreeClassifier()
```

```
#feature scaling
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
xtrain=sc.fit_transform(xtrain)
```

```
xtest=sc.transform(xtest)
```

```
#hypertuning using solver parameter
```

```
logreg=LogisticRegression(solver="saga")
logreg.fit(xtrain,ytrain)
ypred=logreg.predict(xtest)
```

```
print(classification_report(ytest,ypred))
```


	precision	recall	f1-score	support
0	0.96	0.99	0.97	135
1	0.93	0.92	0.93	149
2	0.92	0.89	0.91	168
3	0.93	0.95	0.94	148
accuracy			0.94	600
macro avg	0.94	0.94	0.94	600
weighted avg	0.93	0.94	0.93	600

#checking overfitted on dt

dt.score(xtrain,ytrain)

0.26071428571428573

dt1=DecisionTreeClassifier(max_depth=5)

mymodel(dt1)

	precision	recall	f1-score	support
0	0.95	0.85	0.90	135
1	0.72	0.91	0.81	149
2	0.81	0.71	0.76	168
3	0.88	0.86	0.87	148
accuracy			0.83	600
macro avg	0.84	0.83	0.83	600
weighted avg	0.84	0.83	0.83	600

DecisionTreeClassifier(max_depth=5)

dt2=DecisionTreeClassifier(max_depth=10)

mymodel(dt2)

	precision	recall	f1-score	support
0	0.91	0.88	0.89	135
1	0.77	0.85	0.81	149
2	0.86	0.78	0.82	168
3	0.89	0.92	0.91	148
accuracy			0.85	600
macro avg	0.86	0.86	0.86	600
weighted avg	0.86	0.85	0.86	600

DecisionTreeClassifier(max_depth=10)

for i in range(1,50):

dt1=DecisionTreeClassifier(max_depth=i)

dt2.fit(xtrain,ytrain)

ypred=dt2.predict(xtest)

print(f"{i}= {accuracy_score(ytest,ypred)}")

1= 0.8566666666666667
2= 0.8583333333333333
3= 0.8466666666666667
4= 0.8516666666666667
5= 0.85
6= 0.8483333333333334
7= 0.8666666666666667
8= 0.8466666666666667
9= 0.8516666666666667
10= 0.85
11= 0.8566666666666667
12= 0.865
13= 0.8583333333333333
14= 0.8583333333333333
15= 0.865
16= 0.86
17= 0.8516666666666667
18= 0.8616666666666667
19= 0.8433333333333334
20= 0.8583333333333333
21= 0.86
22= 0.86
23= 0.8533333333333334
24= 0.8533333333333334

```

25= 0.86
26= 0.86
27= 0.8483333333333334
28= 0.855
29= 0.8516666666666667
30= 0.8533333333333334
31= 0.8616666666666667
32= 0.85
33= 0.8566666666666667
34= 0.8616666666666667
35= 0.8533333333333334
36= 0.8583333333333333
37= 0.855
38= 0.8516666666666667
39= 0.855
40= 0.8583333333333333
41= 0.865
42= 0.8533333333333334
43= 0.855
44= 0.855
45= 0.8566666666666667
46= 0.8566666666666667
47= 0.8583333333333333
48= 0.86
49= 0.8566666666666667

```

```
#best value of max_depth=10
```

```
dt3=DecisionTreeClassifier(max_depth=11)
```

```
mymodel(dt3)
```

	precision	recall	f1-score	support
0	0.91	0.85	0.88	135
1	0.75	0.84	0.79	149
2	0.86	0.77	0.81	168
3	0.89	0.95	0.92	148
accuracy			0.85	600
macro avg	0.85	0.85	0.85	600
weighted avg	0.85	0.85	0.85	600

```
DecisionTreeClassifier(max_depth=11)
```

```
#gini
```

```
dt8=DecisionTreeClassifier(criterion="gini",min_samples_leaf=15)
```

```
mymodel(dt8)
```

	precision	recall	f1-score	support
0	0.86	0.85	0.86	135
1	0.75	0.81	0.78	149
2	0.81	0.80	0.80	168
3	0.90	0.86	0.88	148
accuracy			0.83	600
macro avg	0.83	0.83	0.83	600
weighted avg	0.83	0.83	0.83	600

```
DecisionTreeClassifier(min_samples_leaf=15)
```

```
dt9=DecisionTreeClassifier(criterion="gini",max_depth=35)
```

```
mymodel(dt9)
```

	precision	recall	f1-score	support
0	0.93	0.87	0.90	135
1	0.77	0.86	0.81	149
2	0.85	0.77	0.81	168
3	0.88	0.93	0.91	148
accuracy			0.85	600
macro avg	0.86	0.86	0.86	600
weighted avg	0.86	0.85	0.85	600

```
DecisionTreeClassifier(max_depth=35)
```

```
for i in range(1,70):
```

```
dt11=DecisionTreeClassifier(criterion="entropy",max_depth=i)
```

```
dt11.fit(xtrain,ytrain)
```

```

ypred=dt11.predict(xtest)
print(f"{1}={accuracy_score(ytest,ypred)}")
12= 0.8533333333333334
13= 0.8416666666666667
14= 0.8466666666666667
15= 0.835
16= 0.855
17= 0.8416666666666667
18= 0.8466666666666667
19= 0.8466666666666667
20= 0.8433333333333334
21= 0.8466666666666667
22= 0.8316666666666667
23= 0.8433333333333334
24= 0.8566666666666667
25= 0.84
26= 0.835
27= 0.8433333333333334
28= 0.8383333333333334
29= 0.8383333333333334
30= 0.84
31= 0.8483333333333334
32= 0.8366666666666667
33= 0.845
34= 0.8433333333333334
35= 0.8466666666666667
36= 0.8533333333333334
37= 0.8516666666666667
38= 0.8316666666666667
39= 0.8533333333333334
40= 0.8316666666666667
41= 0.8466666666666667
42= 0.84
43= 0.84
44= 0.8433333333333334
45= 0.835
46= 0.83
47= 0.8433333333333334
48= 0.8416666666666667
49= 0.8533333333333334
50= 0.8483333333333334
51= 0.8366666666666667
52= 0.8416666666666667
53= 0.84
54= 0.8533333333333334
55= 0.8516666666666667
56= 0.8516666666666667
57= 0.8466666666666667
58= 0.8416666666666667
59= 0.84
60= 0.8433333333333334
61= 0.845
62= 0.845
63= 0.845
64= 0.8483333333333334
65= 0.8483333333333334
66= 0.85
67= 0.8416666666666667
68= 0.8566666666666667
69= 0.85

```

```

#best value of max_depth when criterion=entropy
dt12=DecisionTreeClassifier(criterion="entropy",max_depth=50)
mymodel(dt12)

```

	precision	recall	f1-score	support
0	0.90	0.84	0.87	135
1	0.75	0.87	0.80	149
2	0.83	0.79	0.81	168
3	0.90	0.87	0.89	148
accuracy			0.84	600
macro avg	0.85	0.84	0.84	600
weighted avg	0.84	0.84	0.84	600

```
DecisionTreeClassifier(criterion='entropy', max_depth=50)
```

```
# import Random Forest classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# instantiate the classifier

rfc = RandomForestClassifier(random_state=0)

# fit the model

rfc.fit(xtrain, ytrain)

# Predict the Test set results

ypred = rfc.predict(xtest)

# Check accuracy score

from sklearn.metrics import accuracy_score

print('Model accuracy score with 10 decision-trees : {0:0.4f}'.format(accuracy_score(ytest, ypred)))
    Model accuracy score with 10 decision-trees : 0.8650

#we import baggingclassifier model
from sklearn.ensemble import BaggingClassifier

bg=BaggingClassifier(LogisticRegression())
bg.fit(xtrain,ytrain)
ypred=bg.predict(xtest)
print(classification_report(ytest,ypred))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	135
1	0.93	0.92	0.92	149
2	0.91	0.89	0.90	168
3	0.93	0.94	0.94	148
accuracy			0.93	600
macro avg	0.93	0.93	0.93	600
weighted avg	0.93	0.93	0.93	600

```
bg=BaggingClassifier(DecisionTreeClassifier())
bg.fit(xtrain,ytrain)
ypred=bg.predict(xtest)
print(classification_report(ytest,ypred))
```

	precision	recall	f1-score	support
0	0.91	0.93	0.92	135
1	0.81	0.83	0.82	149
2	0.81	0.84	0.82	168
3	0.95	0.85	0.90	148
accuracy			0.86	600
macro avg	0.87	0.86	0.86	600
weighted avg	0.86	0.86	0.86	600

```
models=[]
accuracy=[]
models.append(("logistic regression",LogisticRegression()))
models.append(("Decision Tree",DecisionTreeClassifier()))
```

```
models
```

```
[('logistic regression', LogisticRegression()),
 ('Decision Tree', DecisionTreeClassifier())]
```

```
vc=VotingClassifier(estimators=models)
vc.fit(xtrain,ytrain)
ypred=vc.predict(xtest)
print(classification_report(ytest,ypred))
```

	precision	recall	f1-score	support
0	0.89	1.00	0.94	135
1	0.86	0.88	0.87	149
2	0.88	0.84	0.86	168
3	0.96	0.89	0.92	148
accuracy			0.90	600
macro avg	0.90	0.90	0.90	600
weighted avg	0.90	0.90	0.90	600

Boosting

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
gd=GradientBoostingClassifier()
gd.fit(xtrain,ytrain)
ypred=gd.predict(xtest)
print(classification_report(ytest,ypred))
```


	precision	recall	f1-score	support
0	0.97	0.95	0.96	135
1	0.84	0.89	0.87	149
2	0.86	0.83	0.84	168
3	0.93	0.93	0.93	148
accuracy			0.90	600
macro avg	0.90	0.90	0.90	600
weighted avg	0.90	0.90	0.90	600

Here we don't have domain knowledge hence we can not decide the value of K first.
#so we are using elbow method to decide the value of K.

Elbow method

```
wcss=[]
for k in range(1,30):
    km=KMeans(n_clusters=k,init="k-means++",n_init=10,max_iter=300,random_state=1)
    km.fit(x)
    wcss.append(km.inertia_)
```

```
plt.figure(figsize=(5,5))
plt.title("elbow method")
plt.plot(range(1,30),wcss,marker="*")
plt.xlabel("K value")
plt.ylabel("inertia")
plt.show()
```



✓ 0s completed at 13:30

